

Figure 1: Example Set

## Final Project: Neural Style Transfer Quality Improvement

**Xinyuan Peng**  
Columbia University  
xp2177@columbia.edu      **Xinyi Zhang**  
Columbia University  
xz2862@columbia.edu

### Abstract

1 In this project, we'll implement the neural style transfer based on existing study  
 2 by Gatys, etc. Briefly, Neural style transfer will take 3 images, a content image, a  
 3 style image and an input image you want to style, then the style of style image will  
 4 be transferred to the input image with content from content image. Our goal is to  
 5 improve the transfer quality by adjusting intermediate layers used in content and  
 6 style representations and discuss why certain layers are selected instead of others.  
 7 An experimental model will be developed to compare with the baseline model that  
 8 was designed by previous scholars. The performance of the models will be tested  
 9 with respect of training iterations, random seeds and relative weight of content and  
 10 style loss in the loss function.

### 1 Related Work

12 Gatys, etc proposed the artificial system based on deep neural networks in 2015[2], where he discussed  
 13 the structure to build the model based on VGG-19. The method is further explained by the article  
 14 published in the following year[1], where five layers of the content image and one layer of the style  
 15 image are selected from the VGG-19 model. However, it failed to provide evidence to support the  
 16 layer selection. Hotaj[3] tried to visualize the image output of the layers thus support the decision.  
 17 It raised our interest to study features extracted by each layer in the VGG-19 model, and to build a  
 18 more customized style transfer system.

19 Four combinations of content and style are in Figure1.

### 2 Methods

21 First, let's briefly go through the tutorial: Neural Style Transfer with Eager Execution. We reconstruct  
 22 some intermediate layers to get both content and style representations, and we will transform the

23 base input image by minimizing the content and style distance(losses) with back-propagation, where  
24 content loss is sum of the euclidean distance between the two intermediate representations, one comes  
25 from input image and the other one comes from content image; style loss is sum of the mean squared  
26 distance between feature correlation map of the style image and the input image.  
27 The following steps are efforts we made to improve the transfer quality on the basis of the tutorial  
28 with default setting and our own images. The tutorial has introduced how to implement all steps  
29 mentioned in the Gatys et al. paper but there are some differences between the tutorial and the paper  
30 and we will discuss them later.

### 31 **2.1 Replace MaxPooling with AveragePooling**

32 VGG19 in Keras application has MaxPooling2D at the end of each block, but the Gatys et al. paper  
33 mentions average pooling will result into slightly more appealing output than max pooling. After  
34 replacing max pooling with average pooling, we do see slightly differences in detail but not obvious  
35 unless you look at them by pixel. We choose to keep average pooling as faithful to the parper and the  
36 following steps are on the basis of the replacement. The code is available in the appendix.

### 37 **2.2 Overlay Content/Style on Output**

38 Once we have an output image already, we wonder if putting the output image back again(or several  
39 times) as a new content/style image would make the new output better or not. Therefore, we get  
40 three outputs, they are (1)output from average pooling; (2)use (1) as style + original content image;  
41 (3)use (1) as content + original style image. However, we do not find significant differences among 3  
42 outputs. The reason might be (1) does not provides unseen features to (2) and (3).

### 43 **2.3 Change Content Layers**

44 It is worth noting that the tutorial picks `block5_conv2` as `content_layers` but the Gatys et al.  
45 paper chooses to match the content representation on layer `conv4_2`(which is `block4_conv2` in  
46 tensorflow's version of VGG19), we understand that the choice of layers is subjective, but we are  
47 interested in finding out the differences between these two choices, furthermore, comparing with  
48 other applicable choices.

49 As the Gatys et al. paper mentions, we will used the `content_layers` to reconstruct and transform  
50 it into content representation, "lower layers simply reproduce the exact pixel values of original  
51 image", and "higher layers in the network capture the high-level content in terms of objects and  
52 their arrangement in the input image but do not constrain the exact pixel values of reconstruction."  
53 The paper proves higher layers are more appropriate to be `content_layers` but it does not specify  
54 how high these layers should be. Therefore, we compared 5 applicable choices in Figure2, which  
55 shows the outputs of selecting `block3_conv2`, `block3_conv3`, `block4_conv2`, `block4_conv3`,  
56 `block5_conv2` as `content_layers` respectively while holding other parameters all the same. It  
57 is obvious that `block3_conv2` and `block3_conv3` keep too many original pixel value like texture  
58 and color, the trace of style image is that clear. `block5_conv2` shows a strong style but it is hardly  
59 to recognize the object. Therefore, `block4_conv2` and `block4_conv3` may be good choices for  
60 `content_layers` since they balance the trace of styles and the sharpness of objects well.

### 61 **2.4 Add Weights to Style Layers**

62 Previous work of Gatys.ect compute the style loss with equally through five selected layers. To  
63 study the information extracted from the style layers, we adjusted the style weights of the five style  
64 layers. To highlight the effects, we set the weights for 5 selected `style_layers` to 10 respectively,  
65 and keep the weights for other `style_layers` small. For example, in figure3(a), we set the weight  
66 for `block1_conv1` to 10 and set the weight for other layers to 0.2. Since the style representation  
67 computes correlations between different features in different layers, adding weight to a specific layer  
68 will amplify features extract by this layer. The output can be found in figure3.

69 Figure3 shows that `block2_conv1` has the most saturated and warmest color. `block3_conv1`  
70 preserve the silhouette of the bulb and also the street lamps in the back. `block4_conv1` has the most  
71 detailed texture of the style image. `block5_conv1` has the coldest color and highlights the brighter

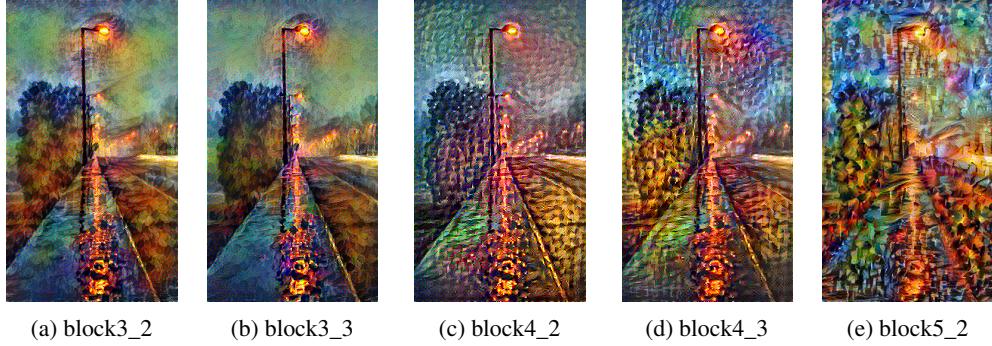


Figure 2: Outputs while Selecting Different Content Layers  
(with  $\alpha/\beta$  set to be  $1e5$ , num\_iteration = 1000, and lr = 5)

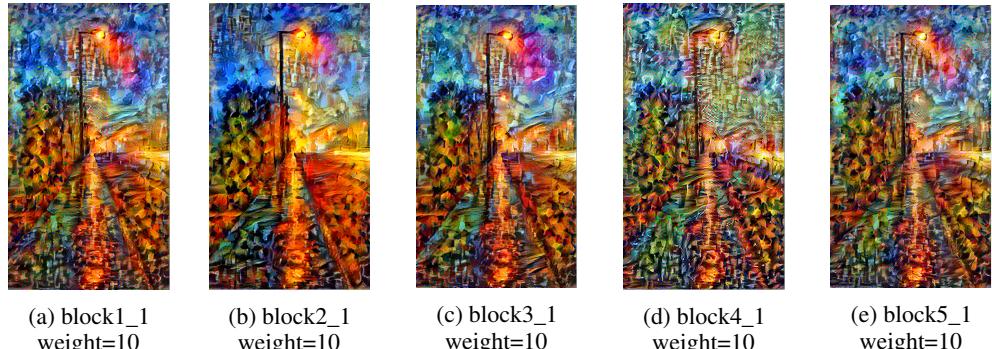


Figure 3: Adjusting weights on style layers  
(with  $\alpha/\beta$  set to be  $1e6$ , num\_iteration = 1000, and lr = 5)

72 color. Compare to the images created by equal loss computation, the adjusted weights pictures tend  
73 to have subtle differences between colors and textures. Again, there's no best output for this model.  
74 The standards are subjective. The code is available in the appendix.

## 75 2.5 Adjust Weights on Content Loss and Style Loss

76 Adjusting the ratio of content loss and style loss during the training process can be informative.  
77 Process focusing on reducing style loss will tend to have output images with significant features and  
78 process focusing on reducing content loss will preserve more flavor of the original content image.  
79 Explained by equation1.

$$L_{total} = \alpha L_{content} + \beta L_{style} \quad (1)$$

80 It's interesting to find out that the style losses we got are around tens of millions while the content  
81 losses are around thousands. To balance the quantity, we start with  $\alpha/\beta = 1e9$  and gradually update  
82 it to  $1e5$ . Outputs are as shown in figure4.

## 83 2.6 Iterations

84 The number of iterations determines the quality of output images together with the learning rate of  
85 the optimizer. For a constant learning rate of 5 in the Adam algorithm, we have the ideal image at  
86 the 200 iterations and there's no further improvement achieved(see Figure5a). However, we would  
87 expect a larger iteration number for a smaller learning rate as the style will transfer slower. We trained  
88 the model with 1000 iterations to maximize the power of parameters. For efficient transformation,  
89 numbers around 200 should be expected.

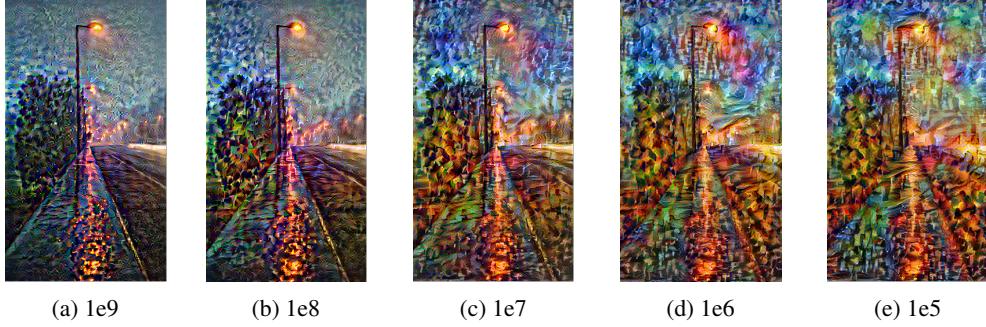


Figure 4: Adjusting weights on content loss and style loss  
(with style layer equally weighted, num\_iteration = 1000, and lr = 5)

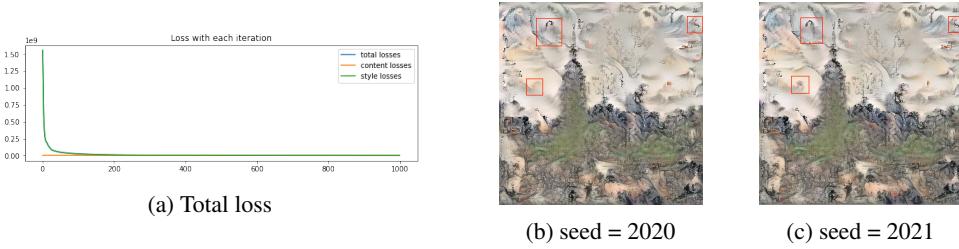


Figure 5: Outputs with  $\alpha/\beta$  set to 1e5, iteration = 1000 and lr = 5

## 90 2.7 Random Seeds

91 Randomness is introduced in the transformation with optimization. Different seeds set before  
92 optimization would lead to output images with subtle differences among colors. To amplify the  
93 effect, we set the ratio of weights on content loss and style loss to be 1e-5. The textures and overall  
94 structures will remain the same. Examples are shown in Figure5b and Figure5c. Different seeds are  
95 set to the two images on the right and marked with red boxes.

## 96 3 Results

97 The final model chooses block4\_conv2 as the content layer, sets weights on different style layers as  
98 (0.05, 0.2, 1, 0.2, 0.2) to keep the structure of the content images with more saturated color. The ratio  
99 of content loss and style loss is set to be 1e3, learning rate equals 5 with 200 iterations. Figure6a  
100 shows the output image of the baseline model while Figure6b shows the final model. Compare to the  
101 baseline model, it is clear to see that we keep the silhouette of the objects in the background with a  
102 clear pattern from the style image. Specifically, the performance of the two models with content1 and  
103 style1 are in Table1.

104 The output images are highly subjective. To have the ideal transformation, following the recommendations  
105 below.

- 106 1. To keep most of the composition of the content image, relatively lower layer (higher than  
107 block3\_conv3) are preferred.
- 108 2. Applying the style to the content image, the ratio between content loss and style loss plays a  
109 big role. Larger ratios tend to keep more painting style of the original content image.

	Running Time	Total Loss	Features
Baseline	111.73	1780934.40	smoother
Final model	23.38	1643762800.00	with clearer structure and more saturated color

Table 1: Model Compare

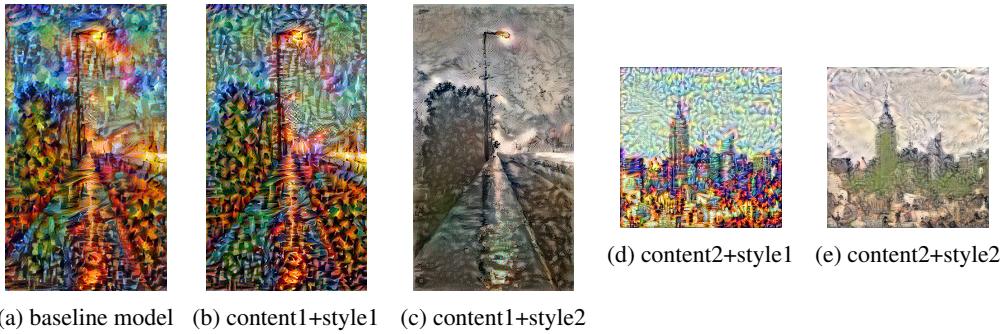


Figure 6: Outputs

- 110        3. To add more detailed textures to the output images, give larger weights to the  
111            block4\_conv1; to add more saturated color, give larger weights to the block2\_conv1.
- 112        4. When efficiency is the priority, larger learning rate with fewer iterations are preferred; when  
113            image quality is more important, train the model with smaller learning rate with more  
114            iterations for fine tuning.

## 115    4 Discussion

116   Neural Style Transfer has a lot of applications that need to be explored. For example, some researchers  
117   are working on use Neural Style Transfer as a data augmentation method, to improve the image  
118   classification accuracy. Our research on improve the transfer quality may provide more specific  
119   features while doing augmentation.

## 120   References

- 121        [1] L. A. Gatys, A. S. Ecker and M. Bethge, "Image Style Transfer Using Convolutional Neural  
122   Networks," 2016 IEEE Conference on Computer Vision and Pattern Recognition (CVPR), Las Vegas,  
123   NV, 2016, pp. 2414-2423, doi: 10.1109/CVPR.2016.265.
- 124        [2] Gatys , Leon A., Alexander S. Ecker, and Matthias Bethge. "A Neural Algorithm of Artistic  
125   Style." arXiv, 2015.
- 126        [3] Hotaj, Eugen. "How to Get Beautiful Results with Neural Style Transfer," March  
127   31, 2020. <https://towardsdatascience.com/how-to-get-beautiful-results-with-neural-style-transfer-75d0c05d6489>.
- 129        [4] Gatys , Leon A. "Google Colaboratory." Google. Google, 2016.  
130   [https://colab.research.google.com/github/tensorflow/models/blob/master/research/nst\\_blogpost/4\\_Neural\\_Style\\_Transfer\\_with\\_Eager\\_Execution.ipynb](https://colab.research.google.com/github/tensorflow/models/blob/master/research/nst_blogpost/4_Neural_Style_Transfer_with_Eager_Execution.ipynb).
- 132        [5] Paul, Sayak. "Implementing Neural Style Transfer Using TensorFlow 2.0," June 21,  
133   2019. <https://www.datacamp.com/community/tutorials/implementing-neural-style-transfer-using-tensorflow>.